

The background features a dark blue grid. A magnifying glass with a silver rim and a dark handle is positioned in the upper right. Two jagged lines, one dark blue and one dark red, trend across the image. The blue line starts high on the left and trends downwards towards the bottom right. The red line starts lower on the left and trends upwards towards the top right, ending under the magnifying glass.

Natural Language Processing for Trend Forecasting

Heládio Sampaio Lopes

Computer Engineering (ITA 2020)

CONTENT



1. INTRODUCTION
2. LITERATURE REVIEW
3. RELATED WORKS
4. MATERIALS AND METHODS
5. ROADMAP



1. INTRODUCTION

2. LITERATURE REVIEW

3. RELATED WORKS

4. MATERIALS AND METHODS

5. ROADMAP

INTRODUCTION



Natural Language Processing can perform a lot of tasks, such as identifying different topics for a set of documents, classifying texts on predefined subjects, and beyond that extract the sentiment to know what people are saying about something.

Motivation

Over the years, more and more knowledge is generated and we humans are not able to process such an amount of information. Natural Language Processing emerges as a technology capable of assisting us in this hard task.

Objectives

Explore Natural Language Processing techniques to propose a framework to modeling in real-time the topics' evolution, and evaluate these models ability to make predictions about future trends.



1. INTRODUCTION

2. LITERATURE REVIEW

3. RELATED WORKS

4. MATERIALS AND METHODS

5. ROADMAP

LITERATURE REVIEW | Text Processing Techniques



The key task to several machine learning problems consist in make a good representation before applying any model. A clean data set can allow a model to increase its performance in the learning process, making a better identification in the patterns present in the variables.

Normalization

Tokenization

Bag of Words

TF-IDF

LITERATURE REVIEW | Text Processing Techniques



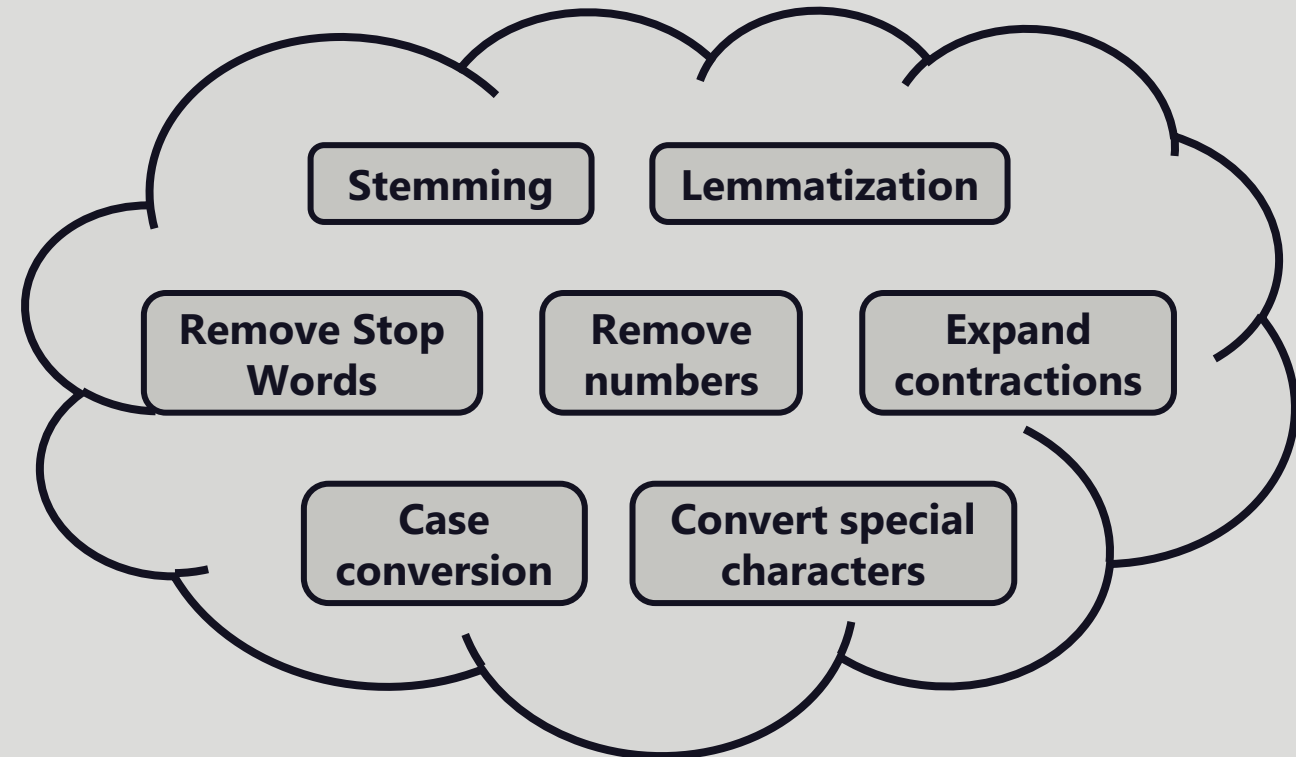
The key task to several machine learning problems consist in make a good representation before applying any model. A clean data set can allow a model to increase its performance in the learning process, making a better identification in the patterns present in the variables.

Normalization

Tokenization

Bag of Words

TF-IDF



LITERATURE REVIEW | Text Processing Techniques



The key task to several machine learning problems consist in make a good representation before applying any model. A clean data set can allow a model to increase its performance in the learning process, making a better identification in the patterns present in the variables.

Normalization

Tokenization

Bag of Words

TF-IDF

0: Something in the way she moves
1: Something in her smile she knows
2: Something in the things she shows me



0: {Something, in, the, way, she, moves}
1: {Something, in, her, smile, she, knows}
2: {Something, in, the, things, she, shows, me}

LITERATURE REVIEW | Text Processing Techniques



The key task to several machine learning problems consist in make a good representation before applying any model. A clean data set can allow a model to increase its performance in the learning process, making a better identification in the patterns present in the variables.

Normalization

Tokenization

Bag of Words

TF-IDF

{ Something, in, the, way, she, moves, her, smile,
knows, things, shows, me }



	Something	In	The	Way	She	Moves	Her	Smile	Knows	Things	Shows	Me
Doc 1:	1	1	1	1	1	1	0	0	0	0	0	0
Doc 2:	1	1	0	0	1	0	1	1	1	0	0	0
Doc 3:	1	1	1	0	1	0	0	0	0	1	1	1

LITERATURE REVIEW | Text Processing Techniques




The key task to several machine learning problems consist in make a good representation before applying any model. A clean data set can allow a model to increase its performance in the learning process, making a better identification in the patterns present in the variables.

Normalization

Tokenization

Bag of Words

TF-IDF



Variables

- $w_{i,j}$: Word Weight
- $\text{tf}_{i,j}$: Term Frequency
- df_i : Document Frequency
- N : Size of Document Set


LITERATURE REVIEW | Word Embedding



The Word Embedding is a technique to represent words in vectors capable of capture the words context in a document. It is also able to smooth the high dimensionality effect by using much more compact vector to represent the words.

Traditional approach

- Word's vector space
- Problems:
 - Sparse representation
 - High dimensionality
 - Distinct vectors for words with the same meaning
 - Have a good day (x) Have a nice day




LITERATURE REVIEW | Word Embedding



The Word Embedding is a technique to represent words in vectors capable of capture the words context in a document. It is also able to smooth the high dimensionality effect by using much more compact vector to represent the words.

Word Embedding


- Dense representation
- Decreases dimensionality
- Similar vectors for words with the same meaning
- Frameworks:
 - Word2Vec
 - Glove
 - Fast Text



LITERATURE REVIEW | Topic Modeling



In NLP, topic modeling is frequently used text-mining tool to identify hidden patterns, called “topics”, in a collection of documents. Latent Dirichlet Allocation, LDA for short, is a statistical model which uses a Dirichlet distribution to model both the topics and the words, after finding the patterns it is human work identify which topics make sense.



Sports



Food



**Sports
Technology**



**Food
Technology**



Technology



CONTENT



1. INTRODUCTION

2. LITERATURE REVIEW

3. RELATED WORKS

4. MATERIALS AND METHODS

5. ROADMAP

RELATED WORKS




Finding meaningful topics in a document collection has been used for a lot of authors for the most various applications like forecast trends. Predicting future trends can be very helpful in various applications, like to model the evolution of research.

Topic Modeling



Jelodar *et al.* (2020) recently use topic modeling in Reddit related posts about the new disease Covid-19 to group similar comments and perform a sentiment analysis.

Trend Forecast



Using pre labeled scholarly articles over 25 years, Shen (2018) performed a neural network forecaster to study the topics growth and codependency between them.




Hurtado *et al.* (2016) use topic modeling to inspect research publications to model the evolution of the direction of research and forecast the near future trends in IT industry.

RELATED WORKS



Hurtado *et al.* (2016) use topic modeling to inspect research publications, patents, and technical reports aiming to model the evolution of the direction of research and forecast the near future trends in IT industry.



Data set containing titles and abstracts over more than six thousand academic papers distributed between 2002 and 2010, data acquired by Tang *et al.* (2008).



Sentence-level association aiming to discover meaningful topics to study the temporal correlation between the topic and predict the popularity of research topics in the future.

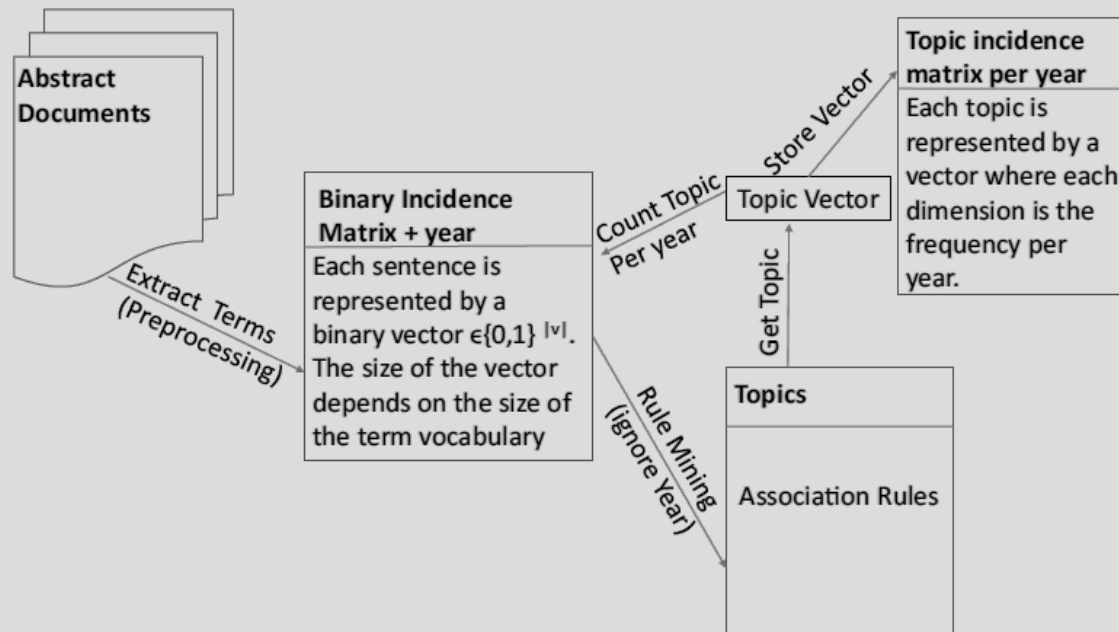


Find topics in the set of articles, model them as an annual time series, forecast those series using the others as input and evaluate the prediction.

RELATED WORKS



Hurtado *et al* (2016) use topic modeling to inspect research publications, patents, and technical reports aiming to model the evolution of the direction of research and forecast the near future trends in IT industry.



- Split the documents in transaction level
- Preprocess the transactions:
 - Basic Normalization
 - Remove common words in scientific publications
- Slight variation of BoW
- Topic discovery and rule refinement
- Topic incidence matrix per year

RELATED WORKS



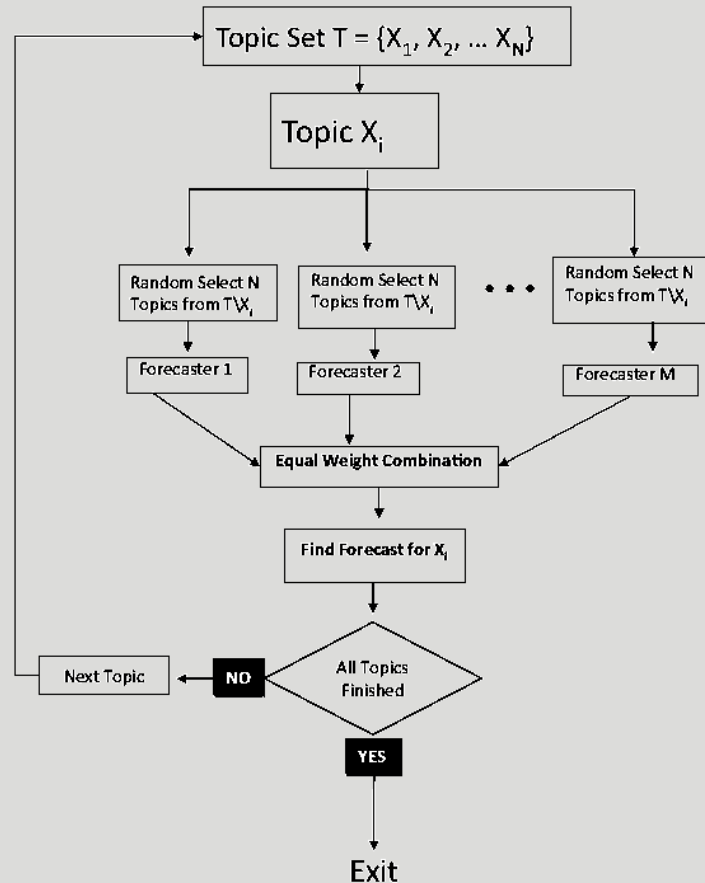
Hurtado *et al*/(2016) use topic modeling to inspect research publications, patents, and technical reports aiming to model the evolution of the direction of research and forecast the near future trends in IT industry.

Topic	2002	2003	2004	2005	2006	2007	2008	2009	2010
Random_walk	0	0	4	4	4	10	9	11	19
Neural_network	14	4	2	2	2	5	7	12	7
Compon_princip	2	3	9	5	12	8	11	7	9
Transfer_learn	0	0	0	0	1	6	10	19	18
Collabor_filter	0	3	10	9	2	8	19	16	20
Select_featur	18	9	28	35	17	32	47	35	77
Topic_model	4	0	5	0	22	27	22	56	36

RELATED WORKS



Hurtado *et al* (2016) use topic modeling to inspect research publications, patents, and technical reports aiming to model the evolution of the direction of research and forecast the near future trends in IT industry.




- Iterate over the topic set
- Ensemble topic forecast
- Predict a topic with others chosen at random
- Average the ensemble
- Evaluate:
 - MSE
 - Accuracy

RELATED WORKS | Gap




Topic discovery and trend forecast were subjects widely explored in the literature. With that in mind, we wish, in this work, to reproduce these techniques. However, in addition to what has been presented we want to be able to explore some modifications.



Real-time system that will keep receiving news in a continuous process.



Redo the discovery process will demand an expensive computational cost.



New topic classification-based system will be proposed in order fill this gap.

CONTENT



1. INTRODUCTION

2. LITERATURE REVIEW

3. RELATED WORKS

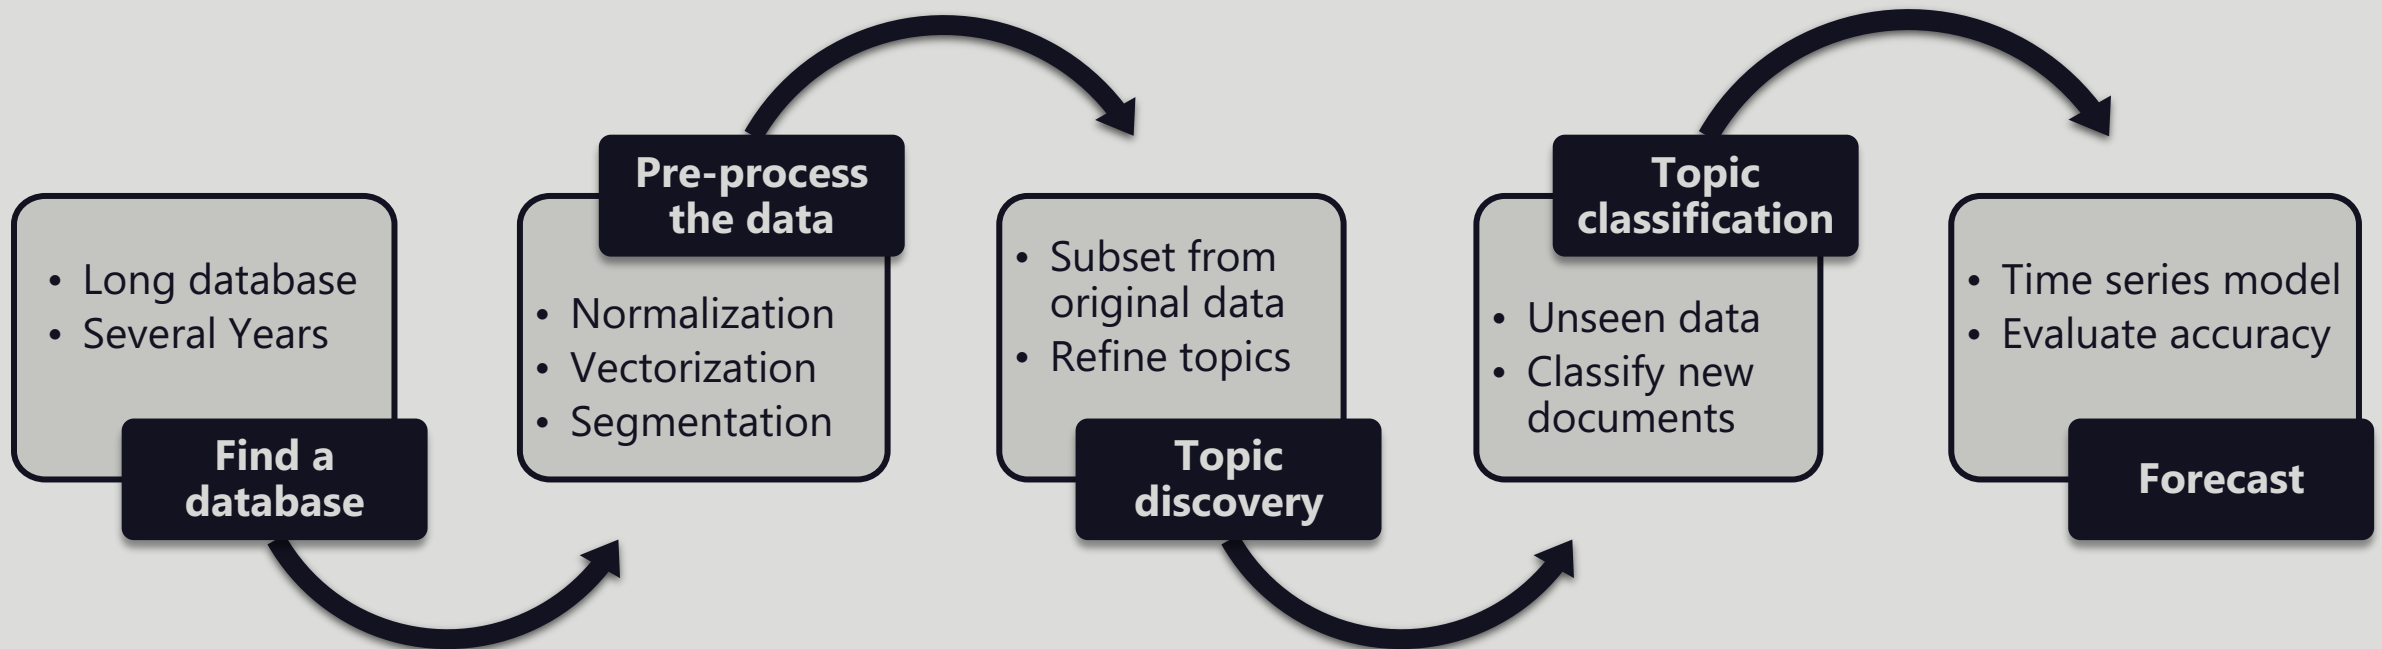
4. MATERIALS AND METHODS

5. ROADMAP

MATERIALS AND METHODS | Objectives



As discussed earlier, we want to build models capable of make predictions regarding the evolution of discovered topics in a set of documents and identify the discovered topics in real time.



MATERIALS AND METHODS | Database



As discussed earlier, we want to build models capable of make predictions regarding the evolution of discovered topics in a set of documents and identify the discovered topics in real time.



Wikipedia Daily News




Newspapers Articles



Academic Papers



Social Media - Reddit




MATERIALS AND METHODS | Pre-processing the Data



As discussed earlier, we want to build models capable of make predictions regarding the evolution of discovered topics in a set of documents and identify the discovered topics in real time.

Task:

Pre-process pipeline to normalize documents, applying the previously seen



T_i




T_m



T_f



Time



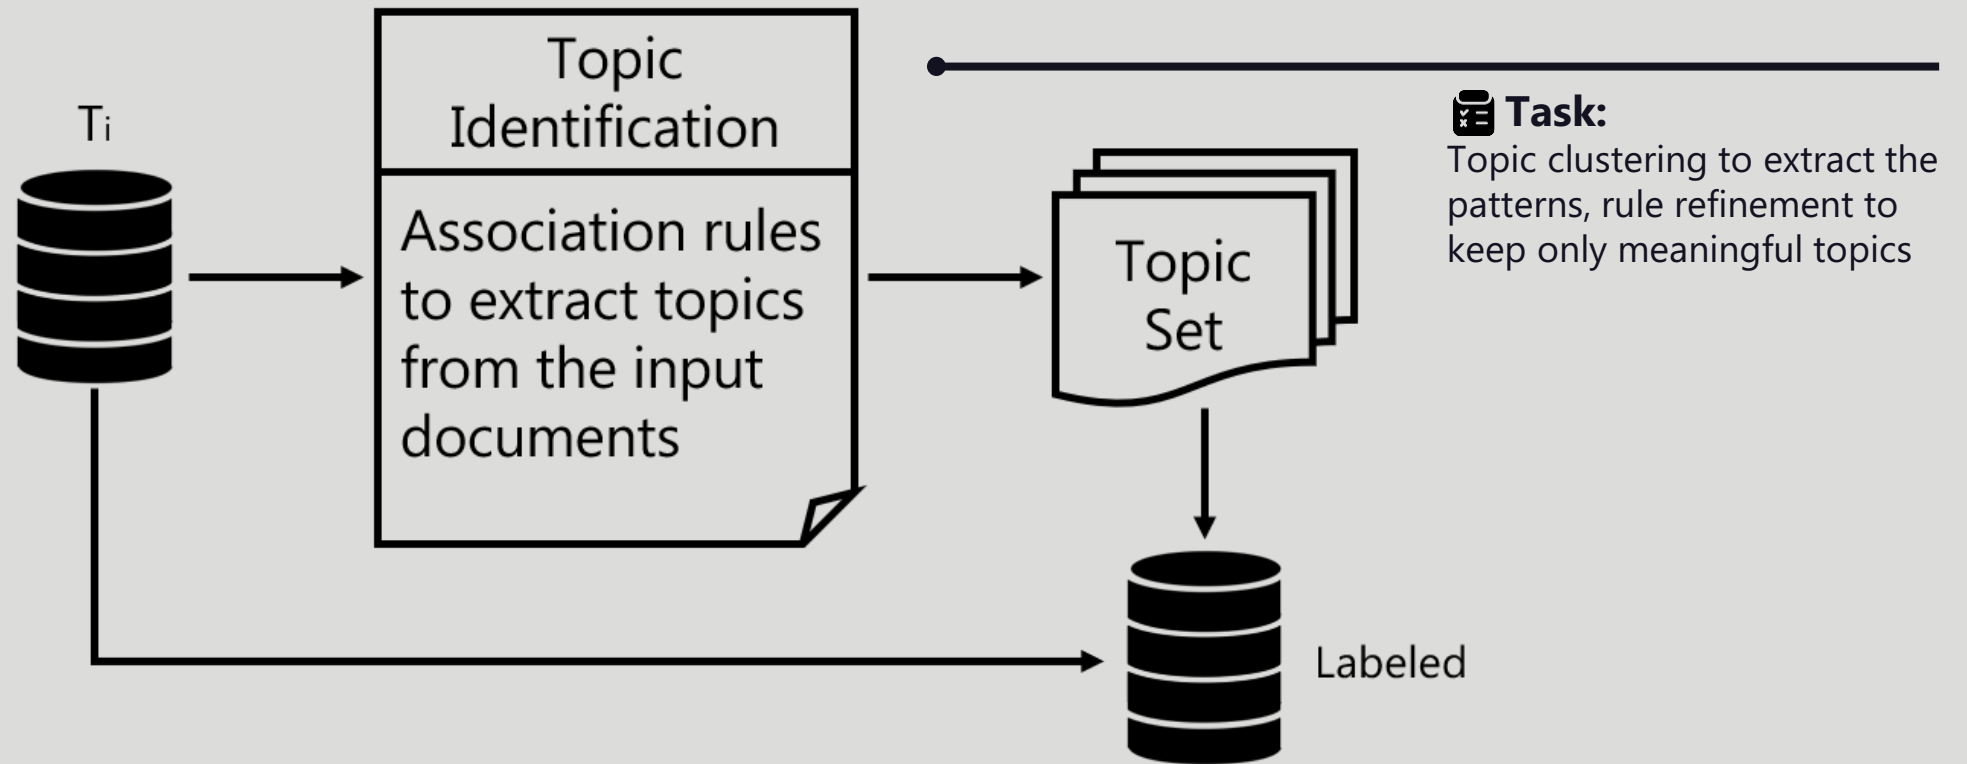
Task:

Database segmentation:
Identifier (t_0-t_j)
Modeler ($t_{j+1}-t_k$)
Validation ($t_{k+1}-t_n$)

MATERIALS AND METHODS | Topic Identification



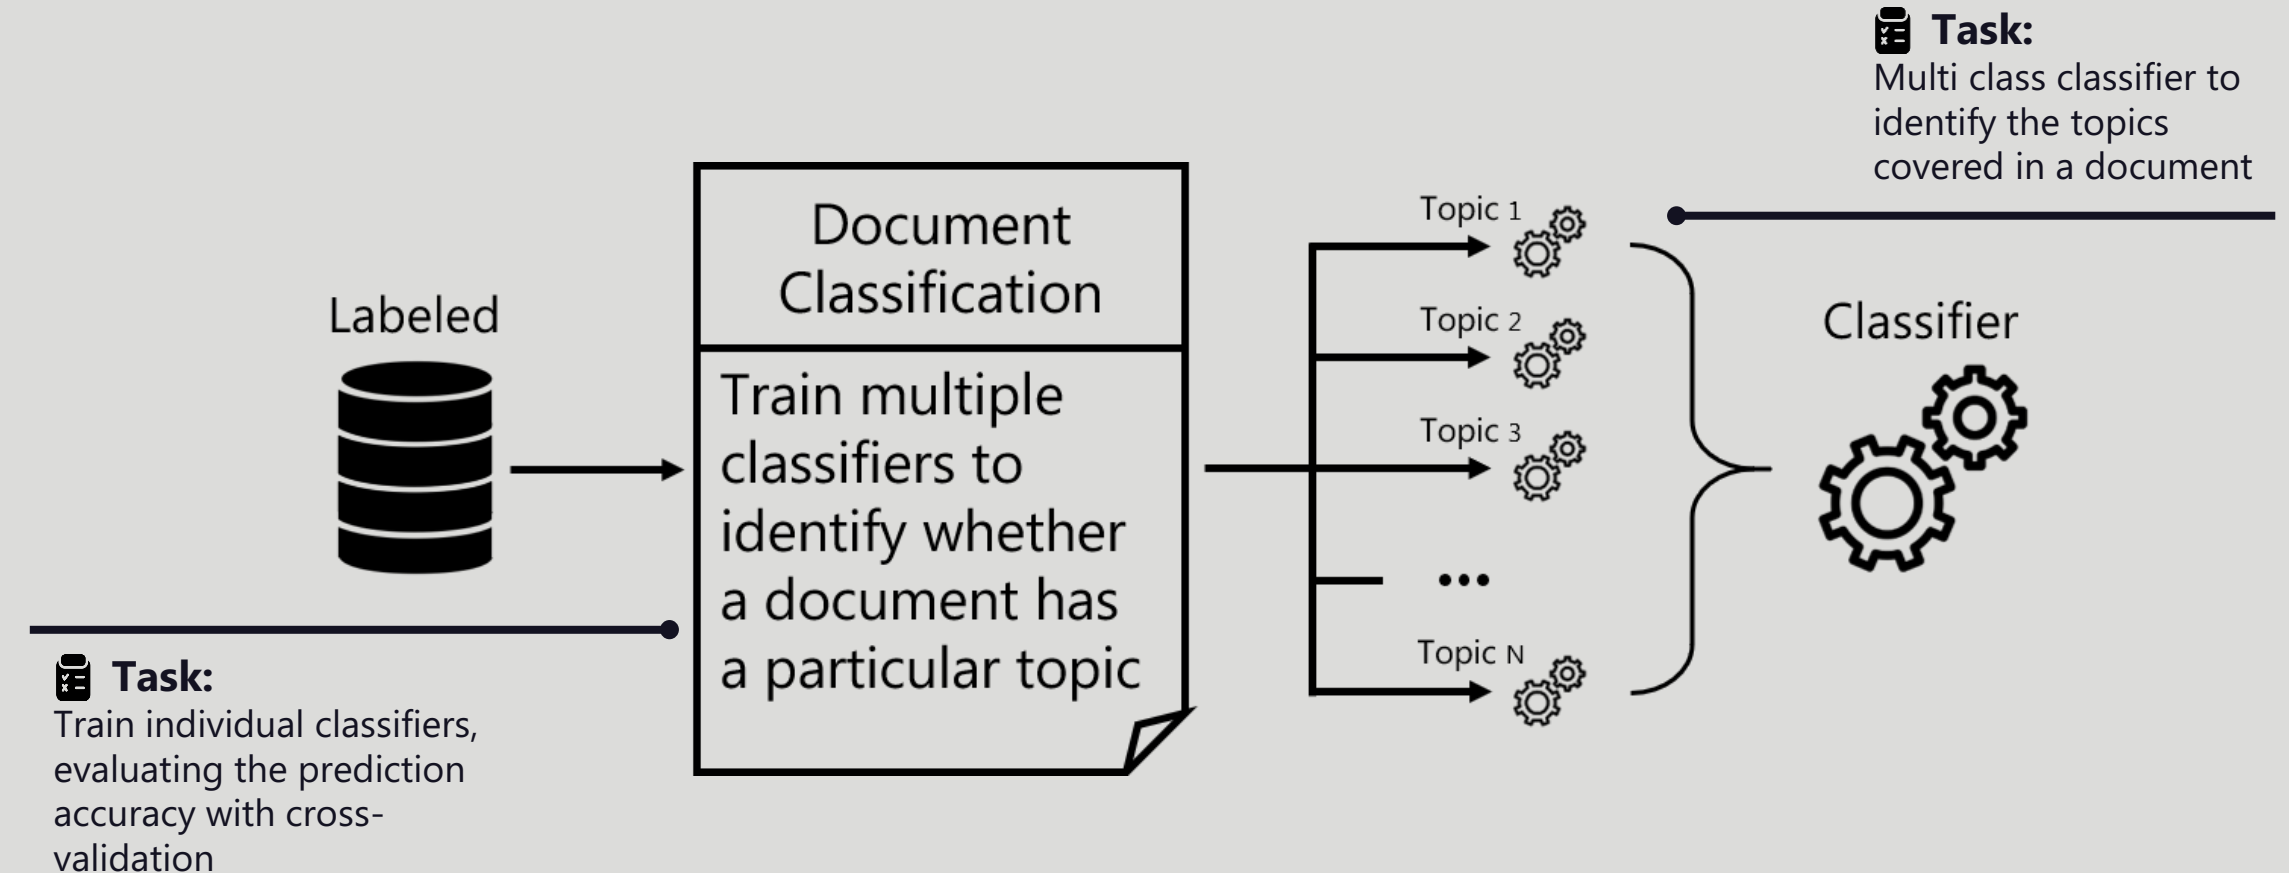
As discussed earlier, we want to build models capable of make predictions regarding the evolution of discovered topics in a set of documents and identify the discovered topics in real time.



MATERIALS AND METHODS | Document Classification



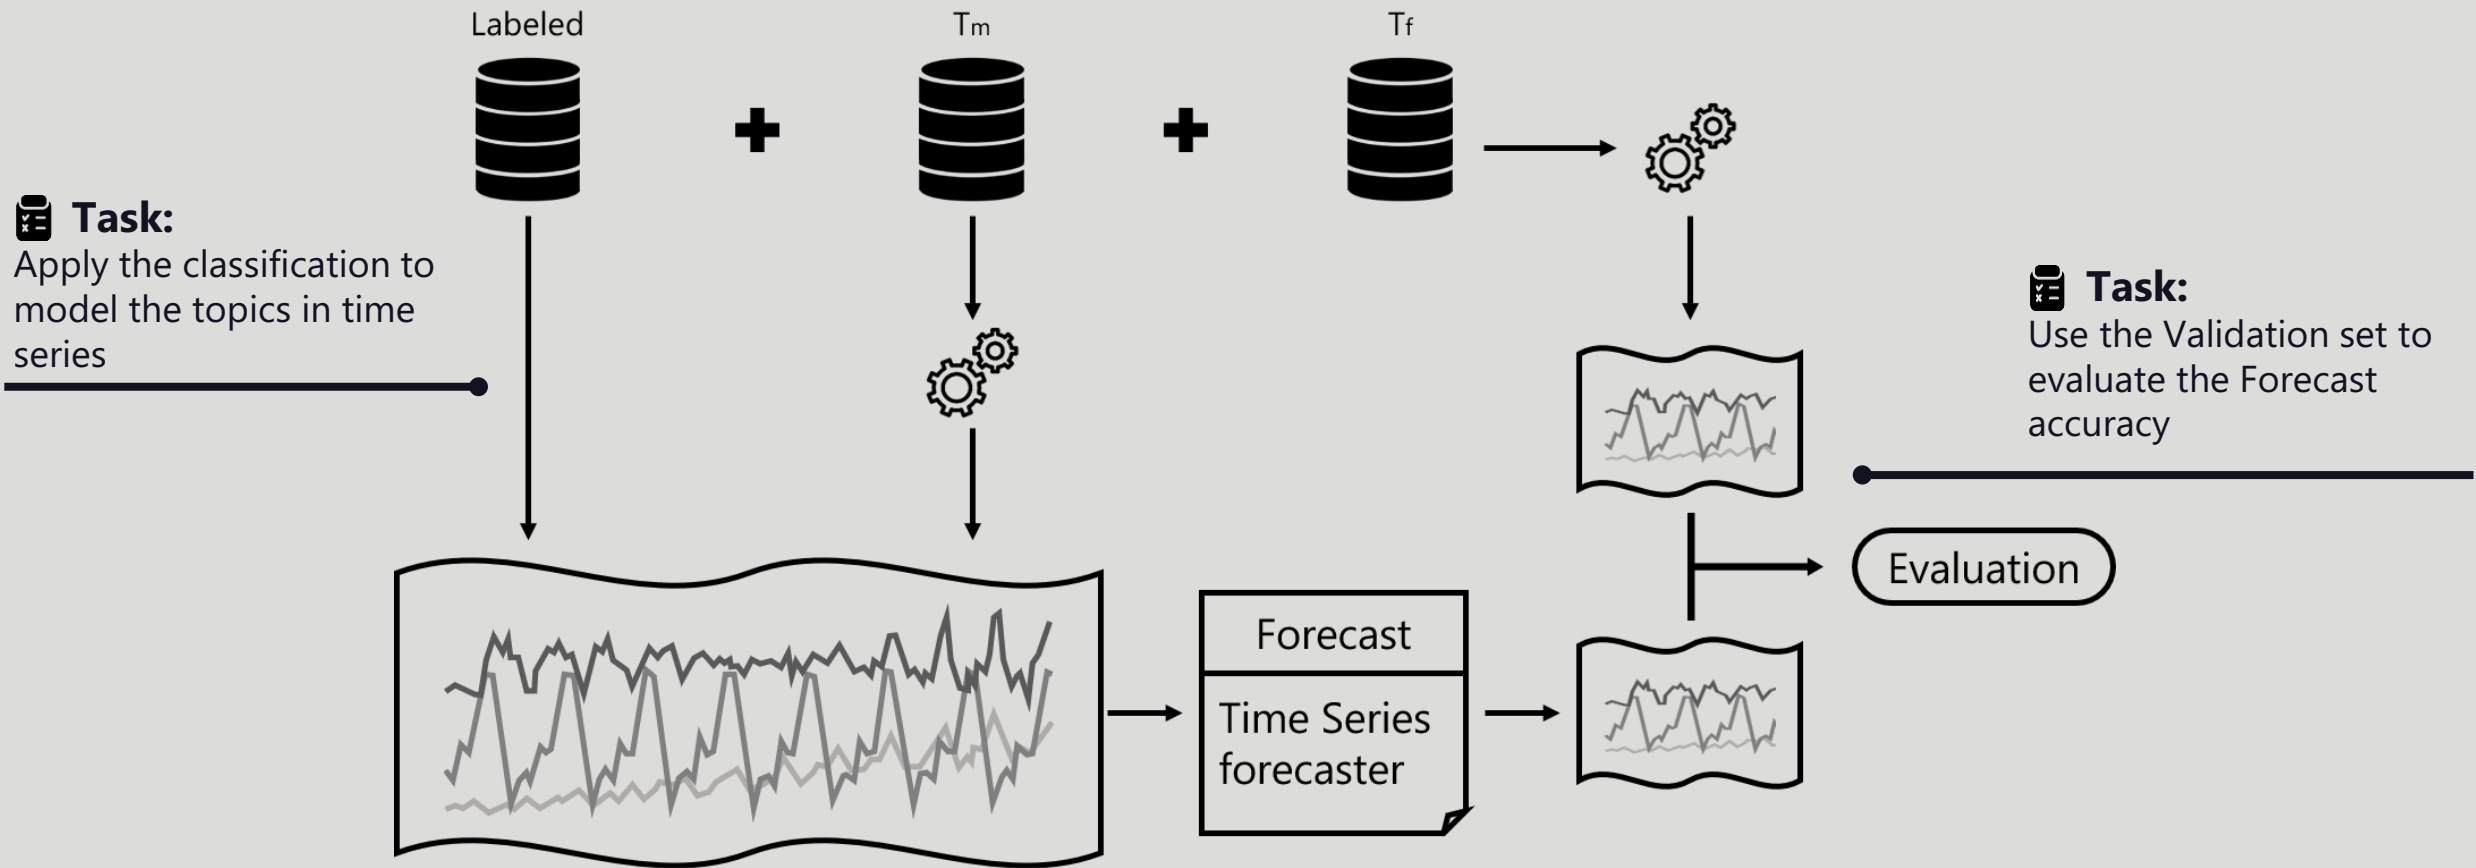
As discussed earlier, we want to build models capable of make predictions regarding the evolution of discovered topics in a set of documents and identify the discovered topics in real time.



MATERIALS AND METHODS | Forecast Evaluation



As discussed earlier, we want to build models capable of make predictions regarding the evolution of discovered topics in a set of documents and identify the discovered topics in real time.



CONTENT



1. INTRODUCTION

2. LITERATURE REVIEW

3. RELATED WORKS

4. MATERIALS AND METHODS

5. ROADMAP

ROADMAP



In view of the problem's complexity, we can elaborate a schedule with the proposed tasks in the previously. The table above show the tasks over the remains months until the end of this work.

Sprint	Start Date	End Date	Duration	Task
#1	August 3	August 16	14 days	- Choose a database - Pre process the database
#2	August 17	September 6	21 days	- Topic Identification
#3	September 7	September 27	21 days	- Document Classification
#4	September 28	October 18	21 days	- Time Series Forecast
#5	October 19	November 8	21 days	- Test and fix bugs

Obrigado!